

## **Lecture 12: Introduction to TensorFlow II**

```
In [2]: # To support both python 2 and python 3
from __future__ import division, print_function, unicode_literals

# Common imports
import numpy as np
import os

# To make this notebook's output stable across runs
def reset_graph(seed=42):
    tf.reset_default_graph()
    tf.set_random_seed(seed)
    np.random.seed(seed)

# To plot pretty figures
%matplotlib inline
import matplotlib
import matplotlib.pyplot as plt
plt.rcParams['axes.labelsize'] = 14
plt.rcParams['xtick.labelsize'] = 12
plt.rcParams['ytick.labelsize'] = 12
```

```
In [3]: import tensorflow as tf
import numpy as np
from sklearn.datasets import fetch_california_housing
reset_graph()

housing = fetch_california_housing()
m, n = housing.data.shape
housing_data_plus_bias = np.c_[np.ones((m, 1)), housing.data]
housing_data_target = housing.target.reshape(-1, 1)
```

```
/Users/mcewen/anaconda3/envs/tensorflow_py35/lib/python3.5/site-packages/h5py/
__init__.py:36: FutureWarning: Conversion of the second argument of issubdtype
from `float` to `np.floating` is deprecated. In future, it will be treated as
`np.float64 == np.dtype(float).type`.
    from ._conv import register_converters as _register_converters
```

```
In [4]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
scaled_housing_data = scaler.fit_transform(housing.data)
scaled_housing_data_plus_bias = np.c_[np.ones((m, 1)), scaled_housing_data]
```

## **Saving and restoring models**

Models and parameters can be saved easily.

Good to not only save the final model but to also checkpoint (i.e. save intermediate models) as training performed.

```
In [6]: with tf.Session() as sess:
        sess.run(init)

        for epoch in range(n_epochs):
            if epoch % 100 == 0:
                print("Epoch", epoch, "MSE =", mse.eval())

                save_path = saver.save(sess, "./my_model.ckpt")
                sess.run(training_op)

        best_theta = theta.eval()
        save_path = saver.save(sess, "./my_model_final.ckpt")
```

```
Epoch 0 MSE = 9.161543
Epoch 100 MSE = 0.7145006
Epoch 200 MSE = 0.56670463
Epoch 300 MSE = 0.5555716
Epoch 400 MSE = 0.5488117
Epoch 500 MSE = 0.5436362
Epoch 600 MSE = 0.53962916
Epoch 700 MSE = 0.53650916
Epoch 800 MSE = 0.5340678
Epoch 900 MSE = 0.53214717
```

```
In [7]: best_theta
```

```
Out[7]: array([[ 2.0685525 ],
               [ 0.8874027 ],
               [ 0.14401658],
               [-0.34770882],
               [ 0.36178368],
               [ 0.00393811],
               [-0.04269556],
               [-0.6614528 ],
               [-0.6375277 ]], dtype=float32)
```

Models can then be restored easily.

```
In [8]: with tf.Session() as sess:  
         saver.restore(sess, "./my_model_final.ckpt")  
         best_theta_restored = theta.eval()
```

```
INFO:tensorflow:Restoring parameters from ./my_model_final.ckpt
```

```
In [9]: best_theta_restored
```

```
Out[9]: array([[ 2.0685525 ],  
               [ 0.8874027 ],  
               [ 0.14401658],  
               [-0.34770882],  
               [ 0.36178368],  
               [ 0.00393811],  
               [-0.04269556],  
               [-0.6614528 ],  
               [-0.6375277 ]], dtype=float32)
```

```
In [10]: np.allclose(best_theta, best_theta_restored)
```

```
Out[10]: True
```



Computational graph definition is saved in file with `.meta` extension.

Can also load graphs.

```
In [11]: reset_graph()  # start with an empty graph.

saver = tf.train.import_meta_graph("./my_model_final.ckpt.meta")  # load the graph
structure
theta = tf.get_default_graph().get_tensor_by_name("theta:0")

with tf.Session() as sess:
    saver.restore(sess, "./my_model_final.ckpt")  # restores the graph's state
    best_theta_restored = theta.eval()
```

```
INFO:tensorflow:Restoring parameters from ./my_model_final.ckpt
```

```
In [12]: np.allclose(best_theta, best_theta_restored)
```

```
Out[12]: True
```

# Visualising computational graphs with TensorBoard

TensorBoard provides functionality to visualise computational graphs and training statistics.

# Logging

```
In [13]: reset_graph()

from datetime import datetime

now = datetime.utcnow().strftime("%Y%m%d%H%M%S")
root_logdir = "tf_logs"
logdir = "{} /run-{}".format(root_logdir, now)
```

```
In [14]: n_epochs = 1000
learning_rate = 0.01

X = tf.placeholder(tf.float32, shape=(None, n + 1), name="X")
y = tf.placeholder(tf.float32, shape=(None, 1), name="y")
theta = tf.Variable(tf.random_uniform([n + 1, 1], -1.0, 1.0, seed=42), name="theta")
y_pred = tf.matmul(X, theta, name="predictions")
error = y_pred - y
mse = tf.reduce_mean(tf.square(error), name="mse")
optimizer = tf.train.GradientDescentOptimizer(learning_rate=learning_rate)
training_op = optimizer.minimize(mse)

init = tf.global_variables_initializer()
```

Create summary and file writer:

```
In [15]: mse_summary = tf.summary.scalar('MSE', mse)
         file_writer = tf.summary.FileWriter(logdir, tf.get_default_graph())
```

```
In [16]: n_epochs = 10
batch_size = 100
n_batches = int(np.ceil(m / batch_size))
```

```
In [17]: def fetch_batch(epoch, batch_index, batch_size):
    np.random.seed(epoch * n_batches + batch_index)
    indices = np.random.randint(m, size=batch_size)
    X_batch = scaled_housing_data_plus_bias[indices]
    y_batch = housing_data_target[indices]
    return X_batch, y_batch
```

```
In [18]: with tf.Session() as sess:
          sess.run(init)

          for epoch in range(n_epochs):
              for batch_index in range(n_batches):
                  X_batch, y_batch = fetch_batch(epoch, batch_index, batch_size)
                  if batch_index % 10 == 0:
                      summary_str = mse_summary.eval(feed_dict={X: X_batch, y: y_batch})
                      step = epoch * n_batches + batch_index
                      file_writer.add_summary(summary_str, step)
                      sess.run(training_op, feed_dict={X: X_batch, y: y_batch})

          best_theta = theta.eval()
```

```
In [19]: file_writer.close()
```

```
In [20]: best_theta
```

```
Out[20]: array([[ 2.0703337 ],
                 [ 0.8637145 ],
                 [ 0.12255151],
                 [-0.31211874],
                 [ 0.38510373],
                 [ 0.00434168],
                 [-0.01232954],
                 [-0.83376896],
                 [-0.8030471 ]], dtype=float32)
```

# TensorBoard

Now we can inspect the logs in TensorBoard.

At the command line run:

```
tensorboard --logdir tf_logs
```

Open in a browser on port 6006, i.e. localhost:6006...

**More complex models**



## Name scopes

Large complex models can easily become cluttered with many nodes making them difficult to visualise directly.

To avoid this problem *name scopes* can be created to group nodes.

```
In [21]: reset_graph()

now = datetime.utcnow().strftime("%Y%m%d%H%M%S")
root_logdir = "tf_logs"
logdir = "{}run-{}".format(root_logdir, now)

n_epochs = 1000
learning_rate = 0.01

X = tf.placeholder(tf.float32, shape=(None, n + 1), name="X")
y = tf.placeholder(tf.float32, shape=(None, 1), name="y")
theta = tf.Variable(tf.random_uniform([n + 1, 1], -1.0, 1.0, seed=42), name="theta")
y_pred = tf.matmul(X, theta, name="predictions")
```

Create error and mse inside loss name scope:

```
In [22]: with tf.name_scope("loss") as scope:  
          error = y_pred - y  
          mse = tf.reduce_mean(tf.square(error), name="mse")
```

Set up optimizer:

```
In [23]: optimizer = tf.train.GradientDescentOptimizer(learning_rate=learning_rate)
training_op = optimizer.minimize(mse)

init = tf.global_variables_initializer()

mse_summary = tf.summary.scalar('MSE', mse)
file_writer = tf.summary.FileWriter(logdir, tf.get_default_graph())
```

Run:

```
In [24]: n_epochs = 10
batch_size = 100
n_batches = int(np.ceil(m / batch_size))

with tf.Session() as sess:
    sess.run(init)

    for epoch in range(n_epochs):
        for batch_index in range(n_batches):
            X_batch, y_batch = fetch_batch(epoch, batch_index, batch_size)
            if batch_index % 10 == 0:
                summary_str = mse_summary.eval(feed_dict={X: X_batch, y: y_batch})
                step = epoch * n_batches + batch_index
                file_writer.add_summary(summary_str, step)
                sess.run(training_op, feed_dict={X: X_batch, y: y_batch})

    best_theta = theta.eval()
```

```
In [25]: file_writer.flush()
file_writer.close()
print("Best theta:")
print(best_theta)
```

Best theta:

```
[[ 2.0703337 ]
 [ 0.8637145 ]
 [ 0.12255151]
 [-0.31211874]
 [ 0.38510373]
 [ 0.00434168]
 [-0.01232954]
 [-0.83376896]
 [-0.8030471 ]]
```

View in TensorBoard...

```
In [26]: print(error.op.name)
```

```
loss/sub
```

```
In [27]: print(mse.op.name)
```

```
loss/mse
```

# Modularity

Many graph components are often repeated.



For example, let's create a graph with two rectified linear units (ReLU).

Could just repeat code but this is not efficient coding and is prone to errors.

```
In [28]: reset_graph()

n_features = 3
X = tf.placeholder(tf.float32, shape=(None, n_features), name="X")

w1 = tf.Variable(tf.random_normal((n_features, 1)), name="weights1")
w2 = tf.Variable(tf.random_normal((n_features, 1)), name="weights2")
b1 = tf.Variable(0.0, name="bias1")
b2 = tf.Variable(0.0, name="bias2")

z1 = tf.add(tf.matmul(X, w1), b1, name="z1")
z2 = tf.add(tf.matmul(X, w2), b2, name="z2")

relu1 = tf.maximum(z1, 0., name="relu1")
relu2 = tf.maximum(z2, 0., name="relu2")

output = tf.add(relu1, relu2, name="output")
```

Better approach is to create a function defining a ReLU.

```
In [29]: reset_graph()

def relu(X):
    w_shape = (int(X.get_shape()[1]), 1)
    w = tf.Variable(tf.random_normal(w_shape), name="weights")
    b = tf.Variable(0.0, name="bias")
    z = tf.add(tf.matmul(X, w), b, name="z")
    return tf.maximum(z, 0., name="relu")

n_features = 3
X = tf.placeholder(tf.float32, shape=(None, n_features), name="X")
relus = [relu(X) for i in range(5)]
output = tf.add_n(relus, name="output")
```

```
In [30]: file_writer = tf.summary.FileWriter("tf_logs/relu1", tf.get_default_graph())
```

View in TensorBoard...

Still somewhat confusing.

Use name scopes to make more clear.

```
In [31]: reset_graph()

def relu(X):
    with tf.name_scope("relu"):
        w_shape = (int(X.get_shape()[1]), 1)
        w = tf.Variable(tf.random_normal(w_shape), name="weights")
        b = tf.Variable(0.0, name="bias")
        z = tf.add(tf.matmul(X, w), b, name="z")
        return tf.maximum(z, 0., name="max")
```

```
In [32]: n_features = 3
X = tf.placeholder(tf.float32, shape=(None, n_features), name="X")
relus = [relu(X) for i in range(5)]
output = tf.add_n(relus, name="output")

file_writer = tf.summary.FileWriter("tf_logs/relu2", tf.get_default_graph())
file_writer.close()
```

# Sharing variables

Often want to share variables between nodes.

For example, say we want to consider ReLUs with non-zero thresholds.

One way is to simply pass parameters around as variables.

```
In [33]: reset_graph()

def relu(X, threshold):
    with tf.name_scope("relu"):
        w_shape = (int(X.get_shape()[1]), 1)
        w = tf.Variable(tf.random_normal(w_shape), name="weights")
        b = tf.Variable(0.0, name="bias")
        z = tf.add(tf.matmul(X, w), b, name="z")
        return tf.maximum(z, threshold, name="max")

threshold = tf.Variable(0.0, name="threshold")
X = tf.placeholder(tf.float32, shape=(None, n_features), name="X")
relus = [relu(X, threshold) for i in range(5)]
output = tf.add_n(relus, name="output")
```

While this works perfectly fine, TensorFlow offers an alternative approach that is cleaner and more modular.

Use TensorFlow `tf.get_variable` function to create (if it doesn't yet exist) or reuse variables.

```
In [34]: reset_graph()

with tf.variable_scope("relu"):
    threshold = tf.get_variable("threshold", shape=(),
                                initializer=tf.constant_initializer(0.0))
```

This will actually raise an exception if the variable has already be created by an earlier call to `tf.get_variable`.

Need to explicitly specify that can reuse variables.

Can be performed by setting variable scope's reuse attribute to True, either by:

```
In [35]: with tf.variable_scope("relu", reuse=True):  
         threshold = tf.get_variable("threshold")
```

or by setting explicitly:

```
In [36]: with tf.variable_scope("relu") as scope:  
         scope.reuse_variables()  
         threshold = tf.get_variable("threshold")
```

Using the shared threshold variable:

```
In [37]: reset_graph()

def relu(X):
    with tf.variable_scope("relu", reuse=True):
        threshold = tf.get_variable("threshold")
        w_shape = int(X.get_shape()[1]), 1
        w = tf.Variable(tf.random_normal(w_shape), name="weights")
        b = tf.Variable(0.0, name="bias")
        z = tf.add(tf.matmul(X, w), b, name="z")
        return tf.maximum(z, threshold, name="max")

X = tf.placeholder(tf.float32, shape=(None, n_features), name="X")
with tf.variable_scope("relu"):
    threshold = tf.get_variable("threshold", shape=(),
                                initializer=tf.constant_initializer(0.0))
relus = [relu(X) for relu_index in range(5)]
output = tf.add_n(relus, name="output")
```

```
In [38]: file_writer = tf.summary.FileWriter("tf_logs/relu3", tf.get_default_graph())
file_writer.close()
```

View in TensorBoard...



Could also put the threshold inside the first ReLU:

```
In [39]: reset_graph()

def relu(X):
    threshold = tf.get_variable("threshold", shape=(),
                                initializer=tf.constant_initializer(0.0))
    w_shape = (int(X.get_shape()[1]), 1)
    w = tf.Variable(tf.random_normal(w_shape), name="weights")
    b = tf.Variable(0.0, name="bias")
    z = tf.add(tf.matmul(X, w), b, name="z")
    return tf.maximum(z, threshold, name="max")

X = tf.placeholder(tf.float32, shape=(None, n_features), name="X")
relus = []
for relu_index in range(5):
    with tf.variable_scope("relu", reuse=(relu_index >= 1)) as scope:
        relus.append(relu(X))
output = tf.add_n(relus, name="output")
```

```
In [40]: file_writer = tf.summary.FileWriter("tf_logs/relu4", tf.get_default_graph())
file_writer.close()
```

View in TensorBoard...

# Neural networks in TensorFlow

## Using TF.Learn (high-level API)

TF.Learn provides a high-level TensorFlow API in Python that is compatible with SciKit-Learn.

```
In [41]: from tensorflow.examples.tutorials.mnist import input_data
mnist = input_data.read_data_sets("/tmp/data/")
```

```
Extracting /tmp/data/train-images-idx3-ubyte.gz
Extracting /tmp/data/train-labels-idx1-ubyte.gz
Extracting /tmp/data/t10k-images-idx3-ubyte.gz
Extracting /tmp/data/t10k-labels-idx1-ubyte.gz
```

```
In [42]: X_train = mnist.train.images
X_test = mnist.test.images
y_train = mnist.train.labels.astype("int")
y_test = mnist.test.labels.astype("int")
```

```
In [43]: import tensorflow as tf

config = tf.contrib.learn.RunConfig(tf_random_seed=42)

feature_cols = tf.contrib.learn.infer_real_valued_columns_from_input(X_train)
dnn_clf = tf.contrib.learn.DNNClassifier(hidden_units=[300,100], n_classes=10,
                                         feature_columns=feature_cols, config=config)
dnn_clf = tf.contrib.learn.SKCompat(dnn_clf)
dnn_clf.fit(X_train, y_train, batch_size=50, steps=40000)
```

```
WARNING:tensorflow:Using temporary folder as model directory: /tmp/tmpe69c9rh9
INFO:tensorflow:Using config: {'_evaluation_master': '', '_keep_checkpoint_every_n_hours': 10000, '_num_ps_replicas': 0, '_log_step_count_steps': 100, '_tf_random_seed': 42, '_task_type': None, '_tf_config': gpu_options {
  per_process_gpu_memory_fraction: 1.0
}, '_save_checkpoints_steps': None, '_environment': 'local', '_cluster_spec': <tensorflow.python.training.server_lib.ClusterSpec object at 0x10abc1f98>, '_master': '', '_save_checkpoints_secs': 600, '_model_dir': '/tmp/tmpe69c9rh9', '_task_id': 0, '_session_config': None, '_save_summary_steps': 100, '_keep_checkpoint_max': 5, '_is_chief': True, '_num_worker_replicas': 0}
WARNING:tensorflow:From /Users/mcewen/anaconda3/envs/tensorflow_py35/lib/python3.5/site-packages/tensorflow/contrib/learn/python/learn/estimators/dnn.py:192: get_global_step (from tensorflow.contrib.framework.python.ops.variables) is deprecated and will be removed in a future version.
Instructions for updating:
Please switch to tf.train.get_global_step
INFO:tensorflow>Create CheckpointSaverHook.
INFO:tensorflow:Saving checkpoints for 1 into /tmp/tmpe69c9rh9/model.ckpt.
INFO:tensorflow:loss = 2.4005778, step = 1
INFO:tensorflow:global_step/sec: 264.728
INFO:tensorflow:loss = 0.31267586, step = 101 (0.379 sec)
INFO:tensorflow:global_step/sec: 234.107
INFO:tensorflow:loss = 0.29636884, step = 201 (0.427 sec)
INFO:tensorflow:global_step/sec: 220.965
```

INFO:tensorflow:loss = 0.4080768, step = 301 (0.452 sec)  
INFO:tensorflow:global\_step/sec: 275.559  
INFO:tensorflow:loss = 0.23435771, step = 401 (0.363 sec)  
INFO:tensorflow:global\_step/sec: 230.905  
INFO:tensorflow:loss = 0.2913079, step = 501 (0.433 sec)  
INFO:tensorflow:global\_step/sec: 162.635  
INFO:tensorflow:loss = 0.07169643, step = 601 (0.615 sec)  
INFO:tensorflow:global\_step/sec: 236.966  
INFO:tensorflow:loss = 0.14009716, step = 701 (0.422 sec)  
INFO:tensorflow:global\_step/sec: 239.115  
INFO:tensorflow:loss = 0.19815028, step = 801 (0.418 sec)  
INFO:tensorflow:global\_step/sec: 246.523  
INFO:tensorflow:loss = 0.11660824, step = 901 (0.406 sec)  
INFO:tensorflow:global\_step/sec: 273.111  
INFO:tensorflow:loss = 0.24129894, step = 1001 (0.366 sec)  
INFO:tensorflow:global\_step/sec: 243.832  
INFO:tensorflow:loss = 0.19023083, step = 1101 (0.410 sec)  
INFO:tensorflow:global\_step/sec: 282.785  
INFO:tensorflow:loss = 0.14697056, step = 1201 (0.353 sec)  
INFO:tensorflow:global\_step/sec: 273.273  
INFO:tensorflow:loss = 0.1890896, step = 1301 (0.366 sec)  
INFO:tensorflow:global\_step/sec: 275.319  
INFO:tensorflow:loss = 0.06954378, step = 1401 (0.364 sec)  
INFO:tensorflow:global\_step/sec: 242.192  
INFO:tensorflow:loss = 0.11753181, step = 1501 (0.413 sec)  
INFO:tensorflow:global\_step/sec: 276.011  
INFO:tensorflow:loss = 0.09583752, step = 1601 (0.362 sec)  
INFO:tensorflow:global\_step/sec: 271.31  
INFO:tensorflow:loss = 0.03823237, step = 1701 (0.369 sec)  
INFO:tensorflow:global\_step/sec: 258.147  
INFO:tensorflow:loss = 0.15002182, step = 1801 (0.387 sec)  
INFO:tensorflow:global\_step/sec: 265.109  
INFO:tensorflow:loss = 0.112414055, step = 1901 (0.377 sec)  
INFO:tensorflow:global\_step/sec: 269.362  
INFO:tensorflow:loss = 0.11133062, step = 2001 (0.371 sec)  
INFO:tensorflow:global\_step/sec: 247.38  
INFO:tensorflow:loss = 0.021828175, step = 2101 (0.404 sec)  
INFO:tensorflow:global\_step/sec: 256.207

INFO:tensorflow:loss = 0.022309918, step = 2201 (0.390 sec)  
INFO:tensorflow:global\_step/sec: 259.115  
INFO:tensorflow:loss = 0.05583959, step = 2301 (0.386 sec)  
INFO:tensorflow:global\_step/sec: 267.849  
INFO:tensorflow:loss = 0.052905086, step = 2401 (0.373 sec)  
INFO:tensorflow:global\_step/sec: 277.507  
INFO:tensorflow:loss = 0.09949044, step = 2501 (0.360 sec)  
INFO:tensorflow:global\_step/sec: 287.202  
INFO:tensorflow:loss = 0.05044488, step = 2601 (0.348 sec)  
INFO:tensorflow:global\_step/sec: 258.488  
INFO:tensorflow:loss = 0.011002336, step = 2701 (0.387 sec)  
INFO:tensorflow:global\_step/sec: 256.5  
INFO:tensorflow:loss = 0.06067799, step = 2801 (0.390 sec)  
INFO:tensorflow:global\_step/sec: 287.312  
INFO:tensorflow:loss = 0.17532235, step = 2901 (0.348 sec)  
INFO:tensorflow:global\_step/sec: 268.433  
INFO:tensorflow:loss = 0.016599359, step = 3001 (0.373 sec)  
INFO:tensorflow:global\_step/sec: 266.434  
INFO:tensorflow:loss = 0.050229833, step = 3101 (0.375 sec)  
INFO:tensorflow:global\_step/sec: 282.718  
INFO:tensorflow:loss = 0.009811628, step = 3201 (0.354 sec)  
INFO:tensorflow:global\_step/sec: 268.309  
INFO:tensorflow:loss = 0.040984496, step = 3301 (0.373 sec)  
INFO:tensorflow:global\_step/sec: 264.772  
INFO:tensorflow:loss = 0.22640564, step = 3401 (0.378 sec)  
INFO:tensorflow:global\_step/sec: 269.333  
INFO:tensorflow:loss = 0.11032523, step = 3501 (0.371 sec)  
INFO:tensorflow:global\_step/sec: 288.665  
INFO:tensorflow:loss = 0.18904746, step = 3601 (0.346 sec)  
INFO:tensorflow:global\_step/sec: 269.212  
INFO:tensorflow:loss = 0.030300684, step = 3701 (0.372 sec)  
INFO:tensorflow:global\_step/sec: 288.321  
INFO:tensorflow:loss = 0.009619746, step = 3801 (0.347 sec)  
INFO:tensorflow:global\_step/sec: 254.618  
INFO:tensorflow:loss = 0.057061974, step = 3901 (0.393 sec)  
INFO:tensorflow:global\_step/sec: 265.451  
INFO:tensorflow:loss = 0.11478825, step = 4001 (0.377 sec)  
INFO:tensorflow:global\_step/sec: 284.962

INFO:tensorflow:loss = 0.030539632, step = 4101 (0.351 sec)  
INFO:tensorflow:global\_step/sec: 284.729  
INFO:tensorflow:loss = 0.054942384, step = 4201 (0.351 sec)  
INFO:tensorflow:global\_step/sec: 259.206  
INFO:tensorflow:loss = 0.18744163, step = 4301 (0.386 sec)  
INFO:tensorflow:global\_step/sec: 266.298  
INFO:tensorflow:loss = 0.1380551, step = 4401 (0.376 sec)  
INFO:tensorflow:global\_step/sec: 286.556  
INFO:tensorflow:loss = 0.0155159235, step = 4501 (0.349 sec)  
INFO:tensorflow:global\_step/sec: 261.465  
INFO:tensorflow:loss = 0.01754779, step = 4601 (0.382 sec)  
INFO:tensorflow:global\_step/sec: 286.819  
INFO:tensorflow:loss = 0.009509798, step = 4701 (0.349 sec)  
INFO:tensorflow:global\_step/sec: 295.747  
INFO:tensorflow:loss = 0.02467858, step = 4801 (0.338 sec)  
INFO:tensorflow:global\_step/sec: 293.854  
INFO:tensorflow:loss = 0.10251294, step = 4901 (0.340 sec)  
INFO:tensorflow:global\_step/sec: 283.519  
INFO:tensorflow:loss = 0.08789965, step = 5001 (0.353 sec)  
INFO:tensorflow:global\_step/sec: 274.044  
INFO:tensorflow:loss = 0.010380711, step = 5101 (0.365 sec)  
INFO:tensorflow:global\_step/sec: 286.238  
INFO:tensorflow:loss = 0.028989911, step = 5201 (0.349 sec)  
INFO:tensorflow:global\_step/sec: 297.931  
INFO:tensorflow:loss = 0.026102195, step = 5301 (0.336 sec)  
INFO:tensorflow:global\_step/sec: 297.045  
INFO:tensorflow:loss = 0.026997259, step = 5401 (0.337 sec)  
INFO:tensorflow:global\_step/sec: 270.202  
INFO:tensorflow:loss = 0.018675137, step = 5501 (0.370 sec)  
INFO:tensorflow:global\_step/sec: 289.067  
INFO:tensorflow:loss = 0.047280237, step = 5601 (0.346 sec)  
INFO:tensorflow:global\_step/sec: 288.248  
INFO:tensorflow:loss = 0.008423564, step = 5701 (0.347 sec)  
INFO:tensorflow:global\_step/sec: 297.896  
INFO:tensorflow:loss = 0.007797351, step = 5801 (0.336 sec)  
INFO:tensorflow:global\_step/sec: 265.633  
INFO:tensorflow:loss = 0.06397591, step = 5901 (0.376 sec)  
INFO:tensorflow:global\_step/sec: 291.179



INFO:tensorflow:loss = 0.090881824, step = 6001 (0.343 sec)  
INFO:tensorflow:global\_step/sec: 282.704  
INFO:tensorflow:loss = 0.014260262, step = 6101 (0.354 sec)  
INFO:tensorflow:global\_step/sec: 295.998  
INFO:tensorflow:loss = 0.016228803, step = 6201 (0.338 sec)  
INFO:tensorflow:global\_step/sec: 288.335  
INFO:tensorflow:loss = 0.048095588, step = 6301 (0.347 sec)  
INFO:tensorflow:global\_step/sec: 284.125  
INFO:tensorflow:loss = 0.040203236, step = 6401 (0.352 sec)  
INFO:tensorflow:global\_step/sec: 294.58  
INFO:tensorflow:loss = 0.012072418, step = 6501 (0.339 sec)  
INFO:tensorflow:global\_step/sec: 288.468  
INFO:tensorflow:loss = 0.013664221, step = 6601 (0.347 sec)  
INFO:tensorflow:global\_step/sec: 250.1  
INFO:tensorflow:loss = 0.01638243, step = 6701 (0.400 sec)  
INFO:tensorflow:global\_step/sec: 293.678  
INFO:tensorflow:loss = 0.008239056, step = 6801 (0.340 sec)  
INFO:tensorflow:global\_step/sec: 295.132  
INFO:tensorflow:loss = 0.012790847, step = 6901 (0.339 sec)  
INFO:tensorflow:global\_step/sec: 283.542  
INFO:tensorflow:loss = 0.021387834, step = 7001 (0.353 sec)  
INFO:tensorflow:global\_step/sec: 295.925  
INFO:tensorflow:loss = 0.008367192, step = 7101 (0.338 sec)  
INFO:tensorflow:global\_step/sec: 290.367  
INFO:tensorflow:loss = 0.04883387, step = 7201 (0.344 sec)  
INFO:tensorflow:global\_step/sec: 298.618  
INFO:tensorflow:loss = 0.009504036, step = 7301 (0.335 sec)  
INFO:tensorflow:global\_step/sec: 291.725  
INFO:tensorflow:loss = 0.016978893, step = 7401 (0.343 sec)  
INFO:tensorflow:global\_step/sec: 286.606  
INFO:tensorflow:loss = 0.010757222, step = 7501 (0.349 sec)  
INFO:tensorflow:global\_step/sec: 291.264  
INFO:tensorflow:loss = 0.010587335, step = 7601 (0.343 sec)  
INFO:tensorflow:global\_step/sec: 293.598  
INFO:tensorflow:loss = 0.005931136, step = 7701 (0.341 sec)  
INFO:tensorflow:global\_step/sec: 296.404  
INFO:tensorflow:loss = 0.005197429, step = 7801 (0.337 sec)  
INFO:tensorflow:global\_step/sec: 293.652

INFO:tensorflow:loss = 0.006449755, step = 7901 (0.341 sec)  
INFO:tensorflow:global\_step/sec: 288.311  
INFO:tensorflow:loss = 0.0020574287, step = 8001 (0.347 sec)  
INFO:tensorflow:global\_step/sec: 295.886  
INFO:tensorflow:loss = 0.006619092, step = 8101 (0.338 sec)  
INFO:tensorflow:global\_step/sec: 295.087  
INFO:tensorflow:loss = 0.03500336, step = 8201 (0.339 sec)  
INFO:tensorflow:global\_step/sec: 290.589  
INFO:tensorflow:loss = 0.025590505, step = 8301 (0.344 sec)  
INFO:tensorflow:global\_step/sec: 292.153  
INFO:tensorflow:loss = 0.005352935, step = 8401 (0.342 sec)  
INFO:tensorflow:global\_step/sec: 290.158  
INFO:tensorflow:loss = 0.0057409974, step = 8501 (0.345 sec)  
INFO:tensorflow:global\_step/sec: 293.607  
INFO:tensorflow:loss = 0.005307211, step = 8601 (0.340 sec)  
INFO:tensorflow:global\_step/sec: 292.729  
INFO:tensorflow:loss = 0.0039272737, step = 8701 (0.342 sec)  
INFO:tensorflow:global\_step/sec: 291.595  
INFO:tensorflow:loss = 0.008358784, step = 8801 (0.343 sec)  
INFO:tensorflow:global\_step/sec: 290.065  
INFO:tensorflow:loss = 0.0027265372, step = 8901 (0.345 sec)  
INFO:tensorflow:global\_step/sec: 241.939  
INFO:tensorflow:loss = 0.0041564466, step = 9001 (0.414 sec)  
INFO:tensorflow:global\_step/sec: 248.57  
INFO:tensorflow:loss = 0.008153919, step = 9101 (0.402 sec)  
INFO:tensorflow:global\_step/sec: 277.052  
INFO:tensorflow:loss = 0.00431604, step = 9201 (0.361 sec)  
INFO:tensorflow:global\_step/sec: 274.892  
INFO:tensorflow:loss = 0.0053035887, step = 9301 (0.364 sec)  
INFO:tensorflow:global\_step/sec: 278.628  
INFO:tensorflow:loss = 0.04991601, step = 9401 (0.359 sec)  
INFO:tensorflow:global\_step/sec: 229.064  
INFO:tensorflow:loss = 0.003466385, step = 9501 (0.436 sec)  
INFO:tensorflow:global\_step/sec: 288.426  
INFO:tensorflow:loss = 0.022571037, step = 9601 (0.347 sec)  
INFO:tensorflow:global\_step/sec: 263.924  
INFO:tensorflow:loss = 0.009801387, step = 9701 (0.380 sec)  
INFO:tensorflow:global\_step/sec: 248.803

INFO:tensorflow:loss = 0.0025435286, step = 9801 (0.400 sec)  
INFO:tensorflow:global\_step/sec: 240.351  
INFO:tensorflow:loss = 0.012849368, step = 9901 (0.416 sec)  
INFO:tensorflow:global\_step/sec: 244.567  
INFO:tensorflow:loss = 0.0025271494, step = 10001 (0.409 sec)  
INFO:tensorflow:global\_step/sec: 245.434  
INFO:tensorflow:loss = 0.0050374116, step = 10101 (0.408 sec)  
INFO:tensorflow:global\_step/sec: 232.404  
INFO:tensorflow:loss = 0.009774216, step = 10201 (0.430 sec)  
INFO:tensorflow:global\_step/sec: 246.513  
INFO:tensorflow:loss = 0.0021551389, step = 10301 (0.406 sec)  
INFO:tensorflow:global\_step/sec: 221.703  
INFO:tensorflow:loss = 0.003974612, step = 10401 (0.450 sec)  
INFO:tensorflow:global\_step/sec: 231.503  
INFO:tensorflow:loss = 0.0045803287, step = 10501 (0.432 sec)  
INFO:tensorflow:global\_step/sec: 244.281  
INFO:tensorflow:loss = 0.010262251, step = 10601 (0.409 sec)  
INFO:tensorflow:global\_step/sec: 256.124  
INFO:tensorflow:loss = 0.026278675, step = 10701 (0.390 sec)  
INFO:tensorflow:global\_step/sec: 269.944  
INFO:tensorflow:loss = 0.004861756, step = 10801 (0.371 sec)  
INFO:tensorflow:global\_step/sec: 264.498  
INFO:tensorflow:loss = 0.0015780836, step = 10901 (0.378 sec)  
INFO:tensorflow:global\_step/sec: 235.682  
INFO:tensorflow:loss = 0.0211305, step = 11001 (0.424 sec)  
INFO:tensorflow:global\_step/sec: 227.749  
INFO:tensorflow:loss = 0.0032920483, step = 11101 (0.439 sec)  
INFO:tensorflow:global\_step/sec: 243.491  
INFO:tensorflow:loss = 0.0009200088, step = 11201 (0.411 sec)  
INFO:tensorflow:global\_step/sec: 251.464  
INFO:tensorflow:loss = 0.0067240377, step = 11301 (0.398 sec)  
INFO:tensorflow:global\_step/sec: 225.963  
INFO:tensorflow:loss = 0.0052556265, step = 11401 (0.443 sec)  
INFO:tensorflow:global\_step/sec: 204.695  
INFO:tensorflow:loss = 0.014270798, step = 11501 (0.488 sec)  
INFO:tensorflow:global\_step/sec: 251.689  
INFO:tensorflow:loss = 0.0010388009, step = 11601 (0.397 sec)  
INFO:tensorflow:global\_step/sec: 268.567

INFO:tensorflow:loss = 0.0033234973, step = 11701 (0.373 sec)  
INFO:tensorflow:global\_step/sec: 238.592  
INFO:tensorflow:loss = 0.0006890038, step = 11801 (0.419 sec)  
INFO:tensorflow:global\_step/sec: 220.398  
INFO:tensorflow:loss = 0.00561754, step = 11901 (0.454 sec)  
INFO:tensorflow:global\_step/sec: 202.065  
INFO:tensorflow:loss = 0.00064186467, step = 12001 (0.494 sec)  
INFO:tensorflow:global\_step/sec: 252.02  
INFO:tensorflow:loss = 0.0026009055, step = 12101 (0.396 sec)  
INFO:tensorflow:global\_step/sec: 263.789  
INFO:tensorflow:loss = 0.0028248266, step = 12201 (0.379 sec)  
INFO:tensorflow:global\_step/sec: 250.232  
INFO:tensorflow:loss = 0.005670745, step = 12301 (0.400 sec)  
INFO:tensorflow:global\_step/sec: 255.048  
INFO:tensorflow:loss = 0.00037431132, step = 12401 (0.392 sec)  
INFO:tensorflow:global\_step/sec: 230.846  
INFO:tensorflow:loss = 0.0009724629, step = 12501 (0.434 sec)  
INFO:tensorflow:global\_step/sec: 246.041  
INFO:tensorflow:loss = 0.002971116, step = 12601 (0.406 sec)  
INFO:tensorflow:global\_step/sec: 238.799  
INFO:tensorflow:loss = 0.0025933585, step = 12701 (0.419 sec)  
INFO:tensorflow:global\_step/sec: 244.936  
INFO:tensorflow:loss = 0.013733688, step = 12801 (0.408 sec)  
INFO:tensorflow:global\_step/sec: 212.487  
INFO:tensorflow:loss = 0.0031792028, step = 12901 (0.471 sec)  
INFO:tensorflow:global\_step/sec: 224.235  
INFO:tensorflow:loss = 0.006552681, step = 13001 (0.446 sec)  
INFO:tensorflow:global\_step/sec: 210.939  
INFO:tensorflow:loss = 0.0034216207, step = 13101 (0.474 sec)  
INFO:tensorflow:global\_step/sec: 244.237  
INFO:tensorflow:loss = 0.0016324178, step = 13201 (0.409 sec)  
INFO:tensorflow:global\_step/sec: 243.963  
INFO:tensorflow:loss = 0.011358909, step = 13301 (0.410 sec)  
INFO:tensorflow:global\_step/sec: 233.174  
INFO:tensorflow:loss = 0.005848825, step = 13401 (0.430 sec)  
INFO:tensorflow:global\_step/sec: 218.895  
INFO:tensorflow:loss = 0.0029444648, step = 13501 (0.457 sec)  
INFO:tensorflow:global\_step/sec: 239.313

INFO:tensorflow:loss = 0.0054142918, step = 13601 (0.417 sec)  
INFO:tensorflow:global\_step/sec: 215.349  
INFO:tensorflow:loss = 0.0021757944, step = 13701 (0.464 sec)  
INFO:tensorflow:global\_step/sec: 259.738  
INFO:tensorflow:loss = 0.0069082426, step = 13801 (0.386 sec)  
INFO:tensorflow:global\_step/sec: 225.309  
INFO:tensorflow:loss = 0.0046300385, step = 13901 (0.443 sec)  
INFO:tensorflow:global\_step/sec: 274.909  
INFO:tensorflow:loss = 0.0021858108, step = 14001 (0.364 sec)  
INFO:tensorflow:global\_step/sec: 253.946  
INFO:tensorflow:loss = 0.011085386, step = 14101 (0.394 sec)  
INFO:tensorflow:global\_step/sec: 257.61  
INFO:tensorflow:loss = 0.006658467, step = 14201 (0.388 sec)  
INFO:tensorflow:global\_step/sec: 265.824  
INFO:tensorflow:loss = 0.0007703595, step = 14301 (0.376 sec)  
INFO:tensorflow:global\_step/sec: 197.294  
INFO:tensorflow:loss = 0.0011539387, step = 14401 (0.508 sec)  
INFO:tensorflow:global\_step/sec: 227.607  
INFO:tensorflow:loss = 0.0010545183, step = 14501 (0.439 sec)  
INFO:tensorflow:global\_step/sec: 216.513  
INFO:tensorflow:loss = 0.0048172306, step = 14601 (0.463 sec)  
INFO:tensorflow:global\_step/sec: 196.147  
INFO:tensorflow:loss = 0.0013384329, step = 14701 (0.509 sec)  
INFO:tensorflow:global\_step/sec: 261.452  
INFO:tensorflow:loss = 0.0011707735, step = 14801 (0.382 sec)  
INFO:tensorflow:global\_step/sec: 263.525  
INFO:tensorflow:loss = 0.002439394, step = 14901 (0.380 sec)  
INFO:tensorflow:global\_step/sec: 245.842  
INFO:tensorflow:loss = 0.0016568727, step = 15001 (0.406 sec)  
INFO:tensorflow:global\_step/sec: 230.473  
INFO:tensorflow:loss = 0.0023113445, step = 15101 (0.434 sec)  
INFO:tensorflow:global\_step/sec: 260.202  
INFO:tensorflow:loss = 0.001015892, step = 15201 (0.384 sec)  
INFO:tensorflow:global\_step/sec: 264.948  
INFO:tensorflow:loss = 0.0022939122, step = 15301 (0.377 sec)  
INFO:tensorflow:global\_step/sec: 186.205  
INFO:tensorflow:loss = 0.0028404044, step = 15401 (0.537 sec)  
INFO:tensorflow:global\_step/sec: 227.866

INFO:tensorflow:loss = 0.005412121, step = 15501 (0.439 sec)  
INFO:tensorflow:global\_step/sec: 230.416  
INFO:tensorflow:loss = 0.0035895878, step = 15601 (0.434 sec)  
INFO:tensorflow:global\_step/sec: 205.284  
INFO:tensorflow:loss = 0.0058891536, step = 15701 (0.487 sec)  
INFO:tensorflow:global\_step/sec: 215.393  
INFO:tensorflow:loss = 0.0008186471, step = 15801 (0.465 sec)  
INFO:tensorflow:global\_step/sec: 241.404  
INFO:tensorflow:loss = 0.00085781404, step = 15901 (0.414 sec)  
INFO:tensorflow:global\_step/sec: 234.486  
INFO:tensorflow:loss = 0.0065113665, step = 16001 (0.426 sec)  
INFO:tensorflow:global\_step/sec: 242.968  
INFO:tensorflow:loss = 0.0030622215, step = 16101 (0.412 sec)  
INFO:tensorflow:global\_step/sec: 209.894  
INFO:tensorflow:loss = 0.0001887983, step = 16201 (0.476 sec)  
INFO:tensorflow:global\_step/sec: 222.774  
INFO:tensorflow:loss = 0.0025533608, step = 16301 (0.450 sec)  
INFO:tensorflow:global\_step/sec: 219.212  
INFO:tensorflow:loss = 0.0011763193, step = 16401 (0.456 sec)  
INFO:tensorflow:global\_step/sec: 227.725  
INFO:tensorflow:loss = 0.0017928132, step = 16501 (0.439 sec)  
INFO:tensorflow:global\_step/sec: 252.458  
INFO:tensorflow:loss = 0.0024241805, step = 16601 (0.396 sec)  
INFO:tensorflow:global\_step/sec: 260.79  
INFO:tensorflow:loss = 0.0024024954, step = 16701 (0.384 sec)  
INFO:tensorflow:global\_step/sec: 252.153  
INFO:tensorflow:loss = 0.002722007, step = 16801 (0.397 sec)  
INFO:tensorflow:global\_step/sec: 170.35  
INFO:tensorflow:loss = 0.001503015, step = 16901 (0.586 sec)  
INFO:tensorflow:global\_step/sec: 224.773  
INFO:tensorflow:loss = 0.003879984, step = 17001 (0.445 sec)  
INFO:tensorflow:global\_step/sec: 241.66  
INFO:tensorflow:loss = 0.0028237433, step = 17101 (0.414 sec)  
INFO:tensorflow:global\_step/sec: 237.813  
INFO:tensorflow:loss = 0.0022598186, step = 17201 (0.421 sec)  
INFO:tensorflow:global\_step/sec: 255.464  
INFO:tensorflow:loss = 0.0011070793, step = 17301 (0.391 sec)  
INFO:tensorflow:global\_step/sec: 233.826

INFO:tensorflow:loss = 0.0013992988, step = 17401 (0.428 sec)  
INFO:tensorflow:global\_step/sec: 220.61  
INFO:tensorflow:loss = 0.0016118982, step = 17501 (0.464 sec)  
INFO:tensorflow:global\_step/sec: 176.551  
INFO:tensorflow:loss = 0.0008287414, step = 17601 (0.556 sec)  
INFO:tensorflow:global\_step/sec: 254.058  
INFO:tensorflow:loss = 0.0010201752, step = 17701 (0.394 sec)  
INFO:tensorflow:global\_step/sec: 242.954  
INFO:tensorflow:loss = 0.00018981178, step = 17801 (0.413 sec)  
INFO:tensorflow:global\_step/sec: 178.664  
INFO:tensorflow:loss = 0.0020044565, step = 17901 (0.560 sec)  
INFO:tensorflow:global\_step/sec: 198.222  
INFO:tensorflow:loss = 0.00046436657, step = 18001 (0.504 sec)  
INFO:tensorflow:global\_step/sec: 172.62  
INFO:tensorflow:loss = 0.00092771056, step = 18101 (0.580 sec)  
INFO:tensorflow:global\_step/sec: 167.684  
INFO:tensorflow:loss = 0.0036528863, step = 18201 (0.596 sec)  
INFO:tensorflow:global\_step/sec: 144.41  
INFO:tensorflow:loss = 0.005607925, step = 18301 (0.692 sec)  
INFO:tensorflow:global\_step/sec: 200.663  
INFO:tensorflow:loss = 0.0026073065, step = 18401 (0.498 sec)  
INFO:tensorflow:global\_step/sec: 178.654  
INFO:tensorflow:loss = 0.0052851615, step = 18501 (0.561 sec)  
INFO:tensorflow:global\_step/sec: 202.498  
INFO:tensorflow:loss = 0.0038767403, step = 18601 (0.494 sec)  
INFO:tensorflow:global\_step/sec: 193.43  
INFO:tensorflow:loss = 0.0017031265, step = 18701 (0.517 sec)  
INFO:tensorflow:global\_step/sec: 207.831  
INFO:tensorflow:loss = 0.0020837628, step = 18801 (0.480 sec)  
INFO:tensorflow:global\_step/sec: 211.753  
INFO:tensorflow:loss = 0.0026109123, step = 18901 (0.473 sec)  
INFO:tensorflow:global\_step/sec: 191.338  
INFO:tensorflow:loss = 0.0014386168, step = 19001 (0.524 sec)  
INFO:tensorflow:global\_step/sec: 209.703  
INFO:tensorflow:loss = 0.0005016105, step = 19101 (0.476 sec)  
INFO:tensorflow:global\_step/sec: 183.617  
INFO:tensorflow:loss = 0.0011798079, step = 19201 (0.545 sec)  
INFO:tensorflow:global\_step/sec: 147.608

INFO:tensorflow:loss = 0.0056989426, step = 19301 (0.677 sec)  
INFO:tensorflow:global\_step/sec: 167.479  
INFO:tensorflow:loss = 0.0014174287, step = 19401 (0.597 sec)  
INFO:tensorflow:global\_step/sec: 195.298  
INFO:tensorflow:loss = 0.00181431, step = 19501 (0.511 sec)  
INFO:tensorflow:global\_step/sec: 263.274  
INFO:tensorflow:loss = 0.00072336383, step = 19601 (0.380 sec)  
INFO:tensorflow:global\_step/sec: 264.501  
INFO:tensorflow:loss = 9.478824e-05, step = 19701 (0.378 sec)  
INFO:tensorflow:global\_step/sec: 269.693  
INFO:tensorflow:loss = 0.00046725423, step = 19801 (0.371 sec)  
INFO:tensorflow:global\_step/sec: 283.992  
INFO:tensorflow:loss = 0.0019611202, step = 19901 (0.352 sec)  
INFO:tensorflow:global\_step/sec: 286.151  
INFO:tensorflow:loss = 0.0025471402, step = 20001 (0.350 sec)  
INFO:tensorflow:global\_step/sec: 285.848  
INFO:tensorflow:loss = 0.00043874516, step = 20101 (0.350 sec)  
INFO:tensorflow:global\_step/sec: 293.607  
INFO:tensorflow:loss = 0.0025391192, step = 20201 (0.341 sec)  
INFO:tensorflow:global\_step/sec: 265.989  
INFO:tensorflow:loss = 0.0010329548, step = 20301 (0.377 sec)  
INFO:tensorflow:global\_step/sec: 295.86  
INFO:tensorflow:loss = 0.00016640488, step = 20401 (0.338 sec)  
INFO:tensorflow:global\_step/sec: 289.613  
INFO:tensorflow:loss = 0.0009935172, step = 20501 (0.345 sec)  
INFO:tensorflow:global\_step/sec: 284.252  
INFO:tensorflow:loss = 0.0008639615, step = 20601 (0.352 sec)  
INFO:tensorflow:global\_step/sec: 283.753  
INFO:tensorflow:loss = 0.00084766647, step = 20701 (0.352 sec)  
INFO:tensorflow:global\_step/sec: 295.259  
INFO:tensorflow:loss = 0.0010883757, step = 20801 (0.339 sec)  
INFO:tensorflow:global\_step/sec: 283.963  
INFO:tensorflow:loss = 0.0015062927, step = 20901 (0.352 sec)  
INFO:tensorflow:global\_step/sec: 295.286  
INFO:tensorflow:loss = 0.0022610624, step = 21001 (0.339 sec)  
INFO:tensorflow:global\_step/sec: 284.517  
INFO:tensorflow:loss = 0.0010078625, step = 21101 (0.352 sec)  
INFO:tensorflow:global\_step/sec: 283.938



INFO:tensorflow:loss = 0.00453955, step = 21201 (0.352 sec)  
INFO:tensorflow:global\_step/sec: 287.58  
INFO:tensorflow:loss = 0.0005150675, step = 21301 (0.348 sec)  
INFO:tensorflow:global\_step/sec: 288.968  
INFO:tensorflow:loss = 0.002108911, step = 21401 (0.346 sec)  
INFO:tensorflow:global\_step/sec: 285.26  
INFO:tensorflow:loss = 0.0003527573, step = 21501 (0.350 sec)  
INFO:tensorflow:global\_step/sec: 292.023  
INFO:tensorflow:loss = 0.0007121074, step = 21601 (0.343 sec)  
INFO:tensorflow:global\_step/sec: 283.626  
INFO:tensorflow:loss = 0.0009194807, step = 21701 (0.353 sec)  
INFO:tensorflow:global\_step/sec: 285.981  
INFO:tensorflow:loss = 0.000383486, step = 21801 (0.349 sec)  
INFO:tensorflow:global\_step/sec: 292.029  
INFO:tensorflow:loss = 0.0004972108, step = 21901 (0.342 sec)  
INFO:tensorflow:global\_step/sec: 284.274  
INFO:tensorflow:loss = 6.2377825e-05, step = 22001 (0.352 sec)  
INFO:tensorflow:global\_step/sec: 277.149  
INFO:tensorflow:loss = 0.0002139326, step = 22101 (0.361 sec)  
INFO:tensorflow:global\_step/sec: 286.635  
INFO:tensorflow:loss = 0.0011587589, step = 22201 (0.349 sec)  
INFO:tensorflow:global\_step/sec: 294.074  
INFO:tensorflow:loss = 0.0015314096, step = 22301 (0.340 sec)  
INFO:tensorflow:global\_step/sec: 285.986  
INFO:tensorflow:loss = 0.0019228004, step = 22401 (0.350 sec)  
INFO:tensorflow:global\_step/sec: 287.245  
INFO:tensorflow:loss = 0.0013467915, step = 22501 (0.348 sec)  
INFO:tensorflow:global\_step/sec: 288.158  
INFO:tensorflow:loss = 0.0023154982, step = 22601 (0.347 sec)  
INFO:tensorflow:global\_step/sec: 288.681  
INFO:tensorflow:loss = 0.00090699946, step = 22701 (0.346 sec)  
INFO:tensorflow:global\_step/sec: 296.74  
INFO:tensorflow:loss = 0.0008416923, step = 22801 (0.337 sec)  
INFO:tensorflow:global\_step/sec: 290.268  
INFO:tensorflow:loss = 0.0020299358, step = 22901 (0.345 sec)  
INFO:tensorflow:global\_step/sec: 285.795  
INFO:tensorflow:loss = 0.00051947305, step = 23001 (0.350 sec)  
INFO:tensorflow:global\_step/sec: 297.033

INFO:tensorflow:loss = 0.0023677696, step = 23101 (0.336 sec)  
INFO:tensorflow:global\_step/sec: 288.996  
INFO:tensorflow:loss = 0.0018610213, step = 23201 (0.346 sec)  
INFO:tensorflow:global\_step/sec: 291.823  
INFO:tensorflow:loss = 0.0018711936, step = 23301 (0.343 sec)  
INFO:tensorflow:global\_step/sec: 293.561  
INFO:tensorflow:loss = 0.0005051533, step = 23401 (0.340 sec)  
INFO:tensorflow:global\_step/sec: 290.037  
INFO:tensorflow:loss = 0.00076665386, step = 23501 (0.345 sec)  
INFO:tensorflow:global\_step/sec: 292.423  
INFO:tensorflow:loss = 0.000556651, step = 23601 (0.342 sec)  
INFO:tensorflow:global\_step/sec: 294.917  
INFO:tensorflow:loss = 0.0001052184, step = 23701 (0.339 sec)  
INFO:tensorflow:global\_step/sec: 293.751  
INFO:tensorflow:loss = 0.00091980194, step = 23801 (0.340 sec)  
INFO:tensorflow:global\_step/sec: 294.418  
INFO:tensorflow:loss = 0.0017980053, step = 23901 (0.340 sec)  
INFO:tensorflow:global\_step/sec: 295.052  
INFO:tensorflow:loss = 0.0014662343, step = 24001 (0.339 sec)  
INFO:tensorflow:global\_step/sec: 294.872  
INFO:tensorflow:loss = 0.0006457205, step = 24101 (0.339 sec)  
INFO:tensorflow:global\_step/sec: 291.595  
INFO:tensorflow:loss = 0.0017085895, step = 24201 (0.343 sec)  
INFO:tensorflow:global\_step/sec: 291.751  
INFO:tensorflow:loss = 0.00019012339, step = 24301 (0.343 sec)  
INFO:tensorflow:global\_step/sec: 293.751  
INFO:tensorflow:loss = 0.0018188378, step = 24401 (0.340 sec)  
INFO:tensorflow:global\_step/sec: 290.217  
INFO:tensorflow:loss = 0.0012226698, step = 24501 (0.345 sec)  
INFO:tensorflow:global\_step/sec: 289.85  
INFO:tensorflow:loss = 0.00072973536, step = 24601 (0.345 sec)  
INFO:tensorflow:global\_step/sec: 296.724  
INFO:tensorflow:loss = 0.0011833684, step = 24701 (0.337 sec)  
INFO:tensorflow:global\_step/sec: 298.927  
INFO:tensorflow:loss = 0.0011878891, step = 24801 (0.334 sec)  
INFO:tensorflow:global\_step/sec: 291.026  
INFO:tensorflow:loss = 0.0013818216, step = 24901 (0.344 sec)  
INFO:tensorflow:global\_step/sec: 290.64

INFO:tensorflow:loss = 0.00048778497, step = 25001 (0.344 sec)  
INFO:tensorflow:global\_step/sec: 290.145  
INFO:tensorflow:loss = 0.00084742083, step = 25101 (0.344 sec)  
INFO:tensorflow:global\_step/sec: 278.89  
INFO:tensorflow:loss = 0.0016339379, step = 25201 (0.359 sec)  
INFO:tensorflow:global\_step/sec: 282.105  
INFO:tensorflow:loss = 0.00015762541, step = 25301 (0.354 sec)  
INFO:tensorflow:global\_step/sec: 284.355  
INFO:tensorflow:loss = 0.00057926326, step = 25401 (0.352 sec)  
INFO:tensorflow:global\_step/sec: 283.233  
INFO:tensorflow:loss = 0.00096197583, step = 25501 (0.353 sec)  
INFO:tensorflow:global\_step/sec: 277.772  
INFO:tensorflow:loss = 0.00072867685, step = 25601 (0.360 sec)  
INFO:tensorflow:global\_step/sec: 280.53  
INFO:tensorflow:loss = 0.0003659475, step = 25701 (0.356 sec)  
INFO:tensorflow:global\_step/sec: 286.732  
INFO:tensorflow:loss = 0.0006396459, step = 25801 (0.349 sec)  
INFO:tensorflow:global\_step/sec: 296.271  
INFO:tensorflow:loss = 0.0014290272, step = 25901 (0.338 sec)  
INFO:tensorflow:global\_step/sec: 295.946  
INFO:tensorflow:loss = 0.00010238056, step = 26001 (0.338 sec)  
INFO:tensorflow:global\_step/sec: 290.226  
INFO:tensorflow:loss = 0.0007521265, step = 26101 (0.345 sec)  
INFO:tensorflow:global\_step/sec: 291.662  
INFO:tensorflow:loss = 0.0011646885, step = 26201 (0.343 sec)  
INFO:tensorflow:global\_step/sec: 291.21  
INFO:tensorflow:loss = 0.0007397069, step = 26301 (0.343 sec)  
INFO:tensorflow:global\_step/sec: 296.281  
INFO:tensorflow:loss = 0.0011290793, step = 26401 (0.338 sec)  
INFO:tensorflow:global\_step/sec: 288.817  
INFO:tensorflow:loss = 0.00079500815, step = 26501 (0.346 sec)  
INFO:tensorflow:global\_step/sec: 291.029  
INFO:tensorflow:loss = 0.00047347747, step = 26601 (0.344 sec)  
INFO:tensorflow:global\_step/sec: 288.322  
INFO:tensorflow:loss = 3.1328407e-05, step = 26701 (0.347 sec)  
INFO:tensorflow:global\_step/sec: 290.106  
INFO:tensorflow:loss = 0.00035584634, step = 26801 (0.345 sec)  
INFO:tensorflow:global\_step/sec: 296.826

INFO:tensorflow:loss = 0.0007996074, step = 26901 (0.337 sec)  
INFO:tensorflow:global\_step/sec: 293.92  
INFO:tensorflow:loss = 0.00077521073, step = 27001 (0.340 sec)  
INFO:tensorflow:global\_step/sec: 293.024  
INFO:tensorflow:loss = 0.00047619123, step = 27101 (0.341 sec)  
INFO:tensorflow:global\_step/sec: 281.852  
INFO:tensorflow:loss = 0.00039861762, step = 27201 (0.355 sec)  
INFO:tensorflow:global\_step/sec: 264.348  
INFO:tensorflow:loss = 0.0011550712, step = 27301 (0.378 sec)  
INFO:tensorflow:global\_step/sec: 250.965  
INFO:tensorflow:loss = 0.0005360307, step = 27401 (0.398 sec)  
INFO:tensorflow:global\_step/sec: 282.026  
INFO:tensorflow:loss = 0.0010285559, step = 27501 (0.355 sec)  
INFO:tensorflow:global\_step/sec: 261.956  
INFO:tensorflow:loss = 0.0011170513, step = 27601 (0.382 sec)  
INFO:tensorflow:global\_step/sec: 281.406  
INFO:tensorflow:loss = 0.00024562303, step = 27701 (0.355 sec)  
INFO:tensorflow:global\_step/sec: 275.104  
INFO:tensorflow:loss = 0.00017662946, step = 27801 (0.363 sec)  
INFO:tensorflow:global\_step/sec: 281.947  
INFO:tensorflow:loss = 0.00060662144, step = 27901 (0.355 sec)  
INFO:tensorflow:global\_step/sec: 230.119  
INFO:tensorflow:loss = 0.0014573914, step = 28001 (0.435 sec)  
INFO:tensorflow:global\_step/sec: 222.396  
INFO:tensorflow:loss = 0.00066597166, step = 28101 (0.451 sec)  
INFO:tensorflow:global\_step/sec: 255.755  
INFO:tensorflow:loss = 0.0007559935, step = 28201 (0.390 sec)  
INFO:tensorflow:global\_step/sec: 291.408  
INFO:tensorflow:loss = 0.00081949023, step = 28301 (0.343 sec)  
INFO:tensorflow:global\_step/sec: 251.939  
INFO:tensorflow:loss = 0.0012622013, step = 28401 (0.397 sec)  
INFO:tensorflow:global\_step/sec: 263.2  
INFO:tensorflow:loss = 0.00015271838, step = 28501 (0.380 sec)  
INFO:tensorflow:global\_step/sec: 270.376  
INFO:tensorflow:loss = 0.0005445875, step = 28601 (0.370 sec)  
INFO:tensorflow:global\_step/sec: 248.966  
INFO:tensorflow:loss = 0.001242336, step = 28701 (0.402 sec)  
INFO:tensorflow:global\_step/sec: 278.546

INFO:tensorflow:loss = 0.0011522959, step = 28801 (0.359 sec)  
INFO:tensorflow:global\_step/sec: 277.429  
INFO:tensorflow:loss = 0.00029501267, step = 28901 (0.361 sec)  
INFO:tensorflow:global\_step/sec: 258.641  
INFO:tensorflow:loss = 0.0012439629, step = 29001 (0.386 sec)  
INFO:tensorflow:global\_step/sec: 271.343  
INFO:tensorflow:loss = 0.0011318299, step = 29101 (0.369 sec)  
INFO:tensorflow:global\_step/sec: 272.929  
INFO:tensorflow:loss = 0.0011576376, step = 29201 (0.366 sec)  
INFO:tensorflow:global\_step/sec: 282.675  
INFO:tensorflow:loss = 0.0012464767, step = 29301 (0.354 sec)  
INFO:tensorflow:global\_step/sec: 213.543  
INFO:tensorflow:loss = 0.0011650472, step = 29401 (0.469 sec)  
INFO:tensorflow:global\_step/sec: 228.147  
INFO:tensorflow:loss = 0.00089854986, step = 29501 (0.437 sec)  
INFO:tensorflow:global\_step/sec: 270.826  
INFO:tensorflow:loss = 0.0003493056, step = 29601 (0.369 sec)  
INFO:tensorflow:global\_step/sec: 284.395  
INFO:tensorflow:loss = 0.0001967982, step = 29701 (0.351 sec)  
INFO:tensorflow:global\_step/sec: 294.581  
INFO:tensorflow:loss = 0.00042677126, step = 29801 (0.339 sec)  
INFO:tensorflow:global\_step/sec: 267.014  
INFO:tensorflow:loss = 0.00096111797, step = 29901 (0.375 sec)  
INFO:tensorflow:global\_step/sec: 280.362  
INFO:tensorflow:loss = 0.00018180076, step = 30001 (0.357 sec)  
INFO:tensorflow:global\_step/sec: 293.301  
INFO:tensorflow:loss = 0.0005855758, step = 30101 (0.341 sec)  
INFO:tensorflow:global\_step/sec: 292.846  
INFO:tensorflow:loss = 0.00021347596, step = 30201 (0.342 sec)  
INFO:tensorflow:global\_step/sec: 282.455  
INFO:tensorflow:loss = 0.00061550434, step = 30301 (0.354 sec)  
INFO:tensorflow:global\_step/sec: 286.051  
INFO:tensorflow:loss = 0.0015167601, step = 30401 (0.349 sec)  
INFO:tensorflow:global\_step/sec: 291.989  
INFO:tensorflow:loss = 0.0007232324, step = 30501 (0.343 sec)  
INFO:tensorflow:global\_step/sec: 297.897  
INFO:tensorflow:loss = 0.00081983523, step = 30601 (0.336 sec)  
INFO:tensorflow:global\_step/sec: 283.147

INFO:tensorflow:loss = 0.0006647414, step = 30701 (0.353 sec)  
INFO:tensorflow:global\_step/sec: 293.197  
INFO:tensorflow:loss = 0.0010431712, step = 30801 (0.341 sec)  
INFO:tensorflow:global\_step/sec: 286.485  
INFO:tensorflow:loss = 0.0005624868, step = 30901 (0.349 sec)  
INFO:tensorflow:global\_step/sec: 287.587  
INFO:tensorflow:loss = 0.00092834665, step = 31001 (0.348 sec)  
INFO:tensorflow:global\_step/sec: 276.7  
INFO:tensorflow:loss = 0.0011234849, step = 31101 (0.361 sec)  
INFO:tensorflow:global\_step/sec: 290.911  
INFO:tensorflow:loss = 0.00031862652, step = 31201 (0.344 sec)  
INFO:tensorflow:global\_step/sec: 293.03  
INFO:tensorflow:loss = 0.0004291339, step = 31301 (0.341 sec)  
INFO:tensorflow:global\_step/sec: 291.187  
INFO:tensorflow:loss = 0.0013877174, step = 31401 (0.343 sec)  
INFO:tensorflow:global\_step/sec: 297.709  
INFO:tensorflow:loss = 0.00020001482, step = 31501 (0.336 sec)  
INFO:tensorflow:global\_step/sec: 299.657  
INFO:tensorflow:loss = 0.00032674294, step = 31601 (0.334 sec)  
INFO:tensorflow:global\_step/sec: 275.276  
INFO:tensorflow:loss = 0.0006365755, step = 31701 (0.364 sec)  
INFO:tensorflow:global\_step/sec: 240.812  
INFO:tensorflow:loss = 0.00016916303, step = 31801 (0.415 sec)  
INFO:tensorflow:global\_step/sec: 290.958  
INFO:tensorflow:loss = 0.0006920603, step = 31901 (0.344 sec)  
INFO:tensorflow:global\_step/sec: 295.674  
INFO:tensorflow:loss = 0.00012926888, step = 32001 (0.338 sec)  
INFO:tensorflow:global\_step/sec: 284.034  
INFO:tensorflow:loss = 0.00039280686, step = 32101 (0.352 sec)  
INFO:tensorflow:global\_step/sec: 289.241  
INFO:tensorflow:loss = 0.0011558913, step = 32201 (0.345 sec)  
INFO:tensorflow:global\_step/sec: 300.043  
INFO:tensorflow:loss = 0.00056561973, step = 32301 (0.333 sec)  
INFO:tensorflow:global\_step/sec: 289.739  
INFO:tensorflow:loss = 0.00017675917, step = 32401 (0.345 sec)  
INFO:tensorflow:global\_step/sec: 289.39  
INFO:tensorflow:loss = 0.00040773995, step = 32501 (0.346 sec)  
INFO:tensorflow:global\_step/sec: 290.326

INFO:tensorflow:loss = 0.00021993961, step = 32601 (0.344 sec)  
INFO:tensorflow:global\_step/sec: 291.508  
INFO:tensorflow:loss = 0.00079787505, step = 32701 (0.343 sec)  
INFO:tensorflow:global\_step/sec: 296.095  
INFO:tensorflow:loss = 0.00090279884, step = 32801 (0.338 sec)  
INFO:tensorflow:global\_step/sec: 293.661  
INFO:tensorflow:loss = 0.0005482138, step = 32901 (0.341 sec)  
INFO:tensorflow:global\_step/sec: 290.595  
INFO:tensorflow:loss = 0.0006554589, step = 33001 (0.344 sec)  
INFO:tensorflow:global\_step/sec: 288.215  
INFO:tensorflow:loss = 0.0018582876, step = 33101 (0.347 sec)  
INFO:tensorflow:global\_step/sec: 294.181  
INFO:tensorflow:loss = 0.00076735055, step = 33201 (0.340 sec)  
INFO:tensorflow:global\_step/sec: 292.27  
INFO:tensorflow:loss = 0.00038163952, step = 33301 (0.342 sec)  
INFO:tensorflow:global\_step/sec: 286.823  
INFO:tensorflow:loss = 0.001371027, step = 33401 (0.349 sec)  
INFO:tensorflow:global\_step/sec: 285.081  
INFO:tensorflow:loss = 0.00029575542, step = 33501 (0.351 sec)  
INFO:tensorflow:global\_step/sec: 280.163  
INFO:tensorflow:loss = 0.00088878634, step = 33601 (0.357 sec)  
INFO:tensorflow:global\_step/sec: 297.187  
INFO:tensorflow:loss = 0.0009937927, step = 33701 (0.336 sec)  
INFO:tensorflow:global\_step/sec: 297.895  
INFO:tensorflow:loss = 0.0005908036, step = 33801 (0.336 sec)  
INFO:tensorflow:global\_step/sec: 289.873  
INFO:tensorflow:loss = 0.0009054421, step = 33901 (0.345 sec)  
INFO:tensorflow:global\_step/sec: 284.707  
INFO:tensorflow:loss = 0.00071158016, step = 34001 (0.351 sec)  
INFO:tensorflow:global\_step/sec: 288.914  
INFO:tensorflow:loss = 0.00053563213, step = 34101 (0.346 sec)  
INFO:tensorflow:global\_step/sec: 293.284  
INFO:tensorflow:loss = 0.0011523266, step = 34201 (0.341 sec)  
INFO:tensorflow:global\_step/sec: 289.047  
INFO:tensorflow:loss = 0.00016720987, step = 34301 (0.346 sec)  
INFO:tensorflow:global\_step/sec: 290.33  
INFO:tensorflow:loss = 0.0004537905, step = 34401 (0.345 sec)  
INFO:tensorflow:global\_step/sec: 292.24

INFO:tensorflow:loss = 0.00044780105, step = 34501 (0.342 sec)  
INFO:tensorflow:global\_step/sec: 288.94  
INFO:tensorflow:loss = 0.00041601766, step = 34601 (0.346 sec)  
INFO:tensorflow:global\_step/sec: 296.694  
INFO:tensorflow:loss = 0.0011050355, step = 34701 (0.337 sec)  
INFO:tensorflow:global\_step/sec: 274.09  
INFO:tensorflow:loss = 0.00052697194, step = 34801 (0.365 sec)  
INFO:tensorflow:global\_step/sec: 185.778  
INFO:tensorflow:loss = 0.00071083516, step = 34901 (0.540 sec)  
INFO:tensorflow:global\_step/sec: 246.85  
INFO:tensorflow:loss = 0.0002585916, step = 35001 (0.403 sec)  
INFO:tensorflow:global\_step/sec: 260.571  
INFO:tensorflow:loss = 0.0008163878, step = 35101 (0.383 sec)  
INFO:tensorflow:global\_step/sec: 190.92  
INFO:tensorflow:loss = 0.00018969526, step = 35201 (0.525 sec)  
INFO:tensorflow:global\_step/sec: 193.444  
INFO:tensorflow:loss = 0.0001349323, step = 35301 (0.517 sec)  
INFO:tensorflow:global\_step/sec: 184.108  
INFO:tensorflow:loss = 0.0010902109, step = 35401 (0.542 sec)  
INFO:tensorflow:global\_step/sec: 203.988  
INFO:tensorflow:loss = 0.00035070805, step = 35501 (0.490 sec)  
INFO:tensorflow:global\_step/sec: 225.803  
INFO:tensorflow:loss = 0.0002819493, step = 35601 (0.443 sec)  
INFO:tensorflow:global\_step/sec: 247.427  
INFO:tensorflow:loss = 0.0008391325, step = 35701 (0.403 sec)  
INFO:tensorflow:global\_step/sec: 294.864  
INFO:tensorflow:loss = 0.000632369, step = 35801 (0.339 sec)  
INFO:tensorflow:global\_step/sec: 289.384  
INFO:tensorflow:loss = 0.00037967376, step = 35901 (0.346 sec)  
INFO:tensorflow:global\_step/sec: 277.089  
INFO:tensorflow:loss = 0.0003320363, step = 36001 (0.361 sec)  
INFO:tensorflow:global\_step/sec: 290.439  
INFO:tensorflow:loss = 0.0011039884, step = 36101 (0.344 sec)  
INFO:tensorflow:global\_step/sec: 277.872  
INFO:tensorflow:loss = 0.0008606437, step = 36201 (0.360 sec)  
INFO:tensorflow:global\_step/sec: 297.916  
INFO:tensorflow:loss = 0.0002927651, step = 36301 (0.336 sec)  
INFO:tensorflow:global\_step/sec: 293.495



INFO:tensorflow:loss = 0.00069891097, step = 36401 (0.341 sec)  
INFO:tensorflow:global\_step/sec: 279.758  
INFO:tensorflow:loss = 0.00015776746, step = 36501 (0.358 sec)  
INFO:tensorflow:global\_step/sec: 271.591  
INFO:tensorflow:loss = 0.00033270434, step = 36601 (0.368 sec)  
INFO:tensorflow:global\_step/sec: 282.735  
INFO:tensorflow:loss = 0.00028696892, step = 36701 (0.354 sec)  
INFO:tensorflow:global\_step/sec: 263.932  
INFO:tensorflow:loss = 0.0008025733, step = 36801 (0.379 sec)  
INFO:tensorflow:global\_step/sec: 284.206  
INFO:tensorflow:loss = 0.0005028858, step = 36901 (0.352 sec)  
INFO:tensorflow:global\_step/sec: 264.347  
INFO:tensorflow:loss = 0.00062359317, step = 37001 (0.378 sec)  
INFO:tensorflow:global\_step/sec: 286.358  
INFO:tensorflow:loss = 0.00019955178, step = 37101 (0.349 sec)  
INFO:tensorflow:global\_step/sec: 296.903  
INFO:tensorflow:loss = 0.00016318852, step = 37201 (0.337 sec)  
INFO:tensorflow:global\_step/sec: 289.063  
INFO:tensorflow:loss = 0.0006694493, step = 37301 (0.346 sec)  
INFO:tensorflow:global\_step/sec: 246.168  
INFO:tensorflow:loss = 0.0003708102, step = 37401 (0.406 sec)  
INFO:tensorflow:global\_step/sec: 245.622  
INFO:tensorflow:loss = 0.00025077042, step = 37501 (0.407 sec)  
INFO:tensorflow:global\_step/sec: 245.825  
INFO:tensorflow:loss = 0.000264987, step = 37601 (0.407 sec)  
INFO:tensorflow:global\_step/sec: 261.187  
INFO:tensorflow:loss = 0.00010086442, step = 37701 (0.383 sec)  
INFO:tensorflow:global\_step/sec: 205.151  
INFO:tensorflow:loss = 0.0006153757, step = 37801 (0.488 sec)  
INFO:tensorflow:global\_step/sec: 224.999  
INFO:tensorflow:loss = 0.0011952921, step = 37901 (0.444 sec)  
INFO:tensorflow:global\_step/sec: 251.428  
INFO:tensorflow:loss = 0.00019914961, step = 38001 (0.398 sec)  
INFO:tensorflow:global\_step/sec: 251.182  
INFO:tensorflow:loss = 0.0008874108, step = 38101 (0.398 sec)  
INFO:tensorflow:global\_step/sec: 246.062  
INFO:tensorflow:loss = 0.00077106326, step = 38201 (0.406 sec)  
INFO:tensorflow:global\_step/sec: 247.689

```
INFO:tensorflow:loss = 5.9410097e-05, step = 38301 (0.404 sec)
INFO:tensorflow:global_step/sec: 244.082
INFO:tensorflow:loss = 0.000106180356, step = 38401 (0.410 sec)
INFO:tensorflow:global_step/sec: 238.792
INFO:tensorflow:loss = 0.0004074581, step = 38501 (0.418 sec)
INFO:tensorflow:global_step/sec: 237.921
INFO:tensorflow:loss = 0.0006691146, step = 38601 (0.422 sec)
INFO:tensorflow:global_step/sec: 230.154
INFO:tensorflow:loss = 0.00045461592, step = 38701 (0.433 sec)
INFO:tensorflow:global_step/sec: 223.848
INFO:tensorflow:loss = 0.00013856331, step = 38801 (0.447 sec)
INFO:tensorflow:global_step/sec: 222.868
INFO:tensorflow:loss = 0.0014249012, step = 38901 (0.449 sec)
INFO:tensorflow:global_step/sec: 237.969
INFO:tensorflow:loss = 0.0001628006, step = 39001 (0.420 sec)
INFO:tensorflow:global_step/sec: 245.218
INFO:tensorflow:loss = 0.0007085502, step = 39101 (0.408 sec)
INFO:tensorflow:global_step/sec: 250.064
INFO:tensorflow:loss = 0.00033850258, step = 39201 (0.400 sec)
INFO:tensorflow:global_step/sec: 225.221
INFO:tensorflow:loss = 0.00038915616, step = 39301 (0.444 sec)
INFO:tensorflow:global_step/sec: 247.8
INFO:tensorflow:loss = 0.0003968734, step = 39401 (0.404 sec)
INFO:tensorflow:global_step/sec: 261.599
INFO:tensorflow:loss = 0.00015025295, step = 39501 (0.382 sec)
INFO:tensorflow:global_step/sec: 245.64
INFO:tensorflow:loss = 0.00062390114, step = 39601 (0.408 sec)
INFO:tensorflow:global_step/sec: 230.621
INFO:tensorflow:loss = 0.00014465627, step = 39701 (0.433 sec)
INFO:tensorflow:global_step/sec: 216.937
INFO:tensorflow:loss = 0.0011388173, step = 39801 (0.461 sec)
INFO:tensorflow:global_step/sec: 232.34
INFO:tensorflow:loss = 0.0007998731, step = 39901 (0.430 sec)
INFO:tensorflow:Saving checkpoints for 40000 into /tmp/tmpe69c9rh9/model.ckpt.
INFO:tensorflow:Loss for final step: 0.00044029002.
```

Out[43]: SKCompat()

```
In [44]: from sklearn.metrics import accuracy_score
```

```
y_pred = dnn_clf.predict(X_test)  
accuracy_score(y_test, y_pred['classes'])
```

```
INFO:tensorflow:Restoring parameters from /tmp/tmpe69c9rh9/model.ckpt-40000
```

```
Out[44]: 0.9835
```

Better accuracy than we achieved with Scitkit-Learn!

## Using plain TensorFlow

The standard API, as we've focused on previously, provides much more control in constructing and training the network architecture.

## Construction of the computational graph

```
In [45]: import tensorflow as tf
```

```
n_inputs = 28*28  # MNIST  
n_hidden1 = 300  
n_hidden2 = 100  
n_outputs = 10
```

```
In [46]: reset_graph()
```

```
X = tf.placeholder(tf.float32, shape=(None, n_inputs), name="X")  
y = tf.placeholder(tf.int64, shape=(None), name="y")
```

Each layer in the network is similar so let's define a general layer that we can reuse.

```
In [47]: def neuron_layer(X, n_neurons, name, activation=None):
    with tf.name_scope(name):
        n_inputs = int(X.get_shape()[1])
        stddev = 2 / np.sqrt(n_inputs) # More on this in next lecture
        init = tf.truncated_normal((n_inputs, n_neurons), stddev=stddev)
        W = tf.Variable(init, name="kernel")
        b = tf.Variable(tf.zeros([n_neurons]), name="bias")
        Z = tf.matmul(X, W) + b
        if activation is not None:
            return activation(Z)
        else:
            return Z
```

Each layer in the network is similar so let's define a general layer that we can reuse.

```
In [47]: def neuron_layer(X, n_neurons, name, activation=None):
    with tf.name_scope(name):
        n_inputs = int(X.get_shape()[1])
        stddev = 2 / np.sqrt(n_inputs) # More on this in next lecture
        init = tf.truncated_normal((n_inputs, n_neurons), stddev=stddev)
        W = tf.Variable(init, name="kernel")
        b = tf.Variable(tf.zeros([n_neurons]), name="bias")
        Z = tf.matmul(X, W) + b
        if activation is not None:
            return activation(Z)
        else:
            return Z
```

In practice, TensorFlow contains many built-in functions so it's generally not necessary to define layers like this (see, e.g., `tf.layers.dense`).

Now let's construct 3 layers.

```
In [48]: with tf.name_scope("dnn"):
          hidden1 = neuron_layer(X, n_hidden1, name="hidden1",
                                activation=tf.nn.relu)
          hidden2 = neuron_layer(hidden1, n_hidden2, name="hidden2",
                                activation=tf.nn.relu)
          logits = neuron_layer(hidden2, n_outputs, name="outputs")
```



Train using cross-entropy, where softmax applied when defining cross-entropy rather than in network construction.

```
In [49]: with tf.name_scope("loss"):
          xentropy = tf.nn.sparse_softmax_cross_entropy_with_logits(labels=y,
                                                                    logits=logits)
          loss = tf.reduce_mean(xentropy, name="loss")
```

Train using cross-entropy, where softmax applied when defining cross-entropy rather than in network construction.

```
In [49]: with tf.name_scope("loss"):
          xentropy = tf.nn.sparse_softmax_cross_entropy_with_logits(labels=y,
                                                                    logits=logits)
          loss = tf.reduce_mean(xentropy, name="loss")
```

Define gradient descent optimizer.

```
In [50]: learning_rate = 0.01

          with tf.name_scope("train"):
              optimizer = tf.train.GradientDescentOptimizer(learning_rate)
              training_op = optimizer.minimize(loss)
```

Define nodes to evaluate accuracy.

The function `in_top_k` ([https://www.tensorflow.org/api\\_docs/python/tf/nn/in\\_top\\_k](https://www.tensorflow.org/api_docs/python/tf/nn/in_top_k)) checks whether the targets (`y`) are in the top `k` predictions (`logits`).

```
In [51]: with tf.name_scope("eval"):  
          correct = tf.nn.in_top_k(logits, y, 1)  
          accuracy = tf.reduce_mean(tf.cast(correct, tf.float32))
```

Define nodes to evaluate accuracy.

The function `in_top_k` ([https://www.tensorflow.org/api\\_docs/python/tf/nn/in\\_top\\_k](https://www.tensorflow.org/api_docs/python/tf/nn/in_top_k)) checks whether the targets (`y`) are in the top `k` predictions (`logits`).

```
In [51]: with tf.name_scope("eval"):  
          correct = tf.nn.in_top_k(logits, y, 1)  
          accuracy = tf.reduce_mean(tf.cast(correct, tf.float32))
```

Finally, define initializer.

```
In [52]: init = tf.global_variables_initializer()  
saver = tf.train.Saver()
```

## Execution (training)

```
In [53]: n_epochs = 10  
batch_size = 50
```

```
In [54]: with tf.Session() as sess:  
    init.run()  
    for epoch in range(n_epochs):  
        for iteration in range(mnist.train.num_examples // batch_size):  
            X_batch, y_batch = mnist.train.next_batch(batch_size)  
            sess.run(training_op, feed_dict={X: X_batch, y: y_batch})  
            acc_train = accuracy.eval(feed_dict={X: X_batch, y: y_batch})  
            acc_val = accuracy.eval(feed_dict={X: mnist.validation.images,  
                                              y: mnist.validation.labels})  
            print(epoch, "Train accuracy:", acc_train, "Val accuracy:", acc_val)  
  
    save_path = saver.save(sess, "./my_model_final.ckpt")
```

```
0 Train accuracy: 0.9 Val accuracy: 0.9146  
1 Train accuracy: 0.94 Val accuracy: 0.9348  
2 Train accuracy: 0.92 Val accuracy: 0.9466  
3 Train accuracy: 0.96 Val accuracy: 0.9508  
4 Train accuracy: 0.92 Val accuracy: 0.9586  
5 Train accuracy: 0.94 Val accuracy: 0.9584  
6 Train accuracy: 0.98 Val accuracy: 0.9608  
7 Train accuracy: 0.96 Val accuracy: 0.9636  
8 Train accuracy: 0.92 Val accuracy: 0.9638  
9 Train accuracy: 0.96 Val accuracy: 0.965
```

## Using the trained network to make predictions

```
In [55]: with tf.Session() as sess:
          saver.restore(sess, "./my_model_final.ckpt")
          X_new_scaled = mnist.test.images[:20]
          Z = logits.eval(feed_dict={X: X_new_scaled})
          y_pred = np.argmax(Z, axis=1)
```

```
INFO:tensorflow:Restoring parameters from ./my_model_final.ckpt
```

Recall that we only defined the network to compute the logits. If require class probabilities then need to apply softmax function. But that is not needed if just want to make single prediction (i.e. just pick class with largest logit value).

```
In [56]: print("Predicted classes:", y_pred)
          print("Actual classes:    ", mnist.test.labels[:20])
```

```
Predicted classes: [7 2 1 0 4 1 4 9 6 9 0 6 9 0 1 5 9 7 3 4]
Actual classes:    [7 2 1 0 4 1 4 9 5 9 0 6 9 0 1 5 9 7 3 4]
```